

ps6

Kehsin Su Esther 3033114294

October 30, 2017

```
knitr::opts_chunk$set(tidy = TRUE, cache = TRUE)
library(pryr)

## Warning: package 'pryr' was built under R version 3.3.3
```

1 Q1

(Brief summary)

The goal of the simulation is to investigate the finite sample properties of the test. How fast does the test statistics(likelihood ratio statistics) converge in distribution to asymptotic distribution and the power. They use power under different samples, proportion, alpha levels to assess their method.

The author has chosen EM algorithm to use maximum likelihood method to obtain 100 sets of parameters. The sample were generated from standard normal distribution. Those values and sample sizes that used to generate null sample may impact the result. Maybe the author should consider more about difference variance of the two components.

The table reveals enough information for the test results. For table 1, personally, I would recommend transfer the row into columns and combine the result of inadjusted test and adjusted test to compare the result more intuitive. I would even calculate the difference of powers. Also, plot some graphs might help reader understand the result better.

When the sample size increase, the it will coverage to alpha. Overall, he results makes more sense when sample size increase and D value decrease.

In my view, more simulation would be more powerful for the hypothesis. Size of 10 might not strong enough. Most of the case, the size should at least be 100 to be representative, and not exceed 1000 for the efficiency. (general 10% of the population)

2 Q2

```
## @knitr database-access

library(RSQLite)
drv <- dbDriver("SQLite")
# relative or absolute path to where the .db file is
dir <- "C:/Users/Esther/Desktop/stat243-fall-2017-master/section/s07"
dbFilename <- "stackoverflow-2016.db"
db <- dbConnect(drv, dbname = file.path(dir, dbFilename))

## @knitr database-tables
```

```

dbListTables(db)

## [1] "answers"          "maxRepByQuestion" "questions"
## [4] "questionsAugment" "questions_tags"   "users"

dbListFields(db, "questions")

## [1] "questionid"      "creationdate" "score"          "viewcount"
## [5] "title"           "ownerid"

dbListFields(db, "users")

## [1] "userid"          "creationdate"   "lastaccessdate" "location"
## [5] "reputation"      "displayname"    "upvotes"         "downvotes"
## [9] "age"             "accountid"

dbListFields(db, "questions_tags")

## [1] "questionid" "tag"

# simple query to get 5 rows from a table
dbGetQuery(db, "select * from questions limit 5")

##   questionid      creationdate score viewcount
## 1   34552550 2016-01-01 00:00:03     0      108
## 2   34552551 2016-01-01 00:00:07     1      151
## 3   34552552 2016-01-01 00:00:39     2     1942
## 4   34552554 2016-01-01 00:00:50     0      153
## 5   34552555 2016-01-01 00:00:51    -1       54
##                                     title
## 1                                     Scope between methods
## 2      Rails - Unknown Attribute - Unable to add a new field to a form on create/update
## 3 Selenium Firefox webdriver won't load a blank page after changing Firefox preferences
## 4                                     Android Studio styles.xml Error
## 5                                     Java: reference to non-final local variables inside a thread
##   ownerid
## 1 5684416
## 2 2457617
## 3 5732525
## 4 5735112
## 5 4646288

dbGetQuery(db, "select * from users limit 5")

##   userid      creationdate      lastaccessdate
## 1    13 2008-08-01 04:18:04 2017-03-13 21:03:28
## 2    24 2008-08-01 12:12:53 2017-03-10 22:11:55
## 3    36 2008-08-01 12:43:55 2017-03-13 20:33:22
## 4    37 2008-08-01 12:44:00 2017-03-13 18:08:20
## 5    39 2008-08-01 12:44:55 2017-03-11 13:17:47
##                                     location reputation      displayname upvotes
## 1 Raleigh, NC, United States      157267 Chris Jester-Young      5198
## 2   London, United Kingdom        2321      sanmiguel          535
## 3      San Diego, CA              4038          Pat            423
## 4   Georgetown, Canada           4269      Wally Lawless        257
## 5    Istanbul, Turkey            10172      huseyint            523
##   downvotes age accountid
## 1      210  37          9
## 2         3  34         17
## 3        10  36         27
## 4        18  36         28
## 5        21  32         30

```

```

dbGetQuery(db, "select * from questions_tags limit 5")

##      questionid      tag
## 1    34552711      c#
## 2    34552711    razor
## 3    34552711    flags
## 4    34552829 javascript
## 5    34552829    rxjs

# select the unique userid that ask questions contain both python and r,
# then exclude those contain python.
result <- dbGetQuery(db, "select distinct userid from
    questions Q
    join questions_tags T on Q.questionid = T.questionid
    join users U on Q.ownerid = U.userid
    where tag = 'r' or tag = 'python'
    except
    select distinct userid from
    questions Q
    join questions_tags T on Q.questionid = T.questionid
    join users U on Q.ownerid = U.userid
    where tag = 'python'
    ")
# summary the result
dim(result)

## [1] 18611      1

```

3 Q3

The following codes are refer from chapter 8.

```

#from pyspark import SparkFiles
dir = '/global/scratch/paciorek/wikistats_full'
#dir=$HOME
#lines = sc.textFile(dir+'/' + 'newdir')
lines = sc.textFile(dir + '/' + 'dated')

lines.getNumPartitions() # 16590 (480 input files) for full dataset

# note delayed evaluation
lines.count()
testLines = lines.take(10)
testLines[0]
testLines[9]

### filter to sites of interest ###

import re
from operator import add

def find(line, regex = "Disney", language = None):
    vals = line.split(' ')
    if len(vals) < 6:
        return(False)

```

```

    tmp = re.search(regex, vals[3])
    if tmp is None or (language != None and vals[2] != language):
        return(False)
    else:
        return(True)

lines.filter(find).take(100) # pretty quick

# not clear if should repartition; will likely have small partitions if not
obama = lines.filter(find).repartition(480) # ~ 18 minutes for full dataset (but remember lazy evaluation)
obama.count() # 433k observations for full dataset
mydir = '/global/scratch/esther730'
outputDir = mydir+'/' + 'Dinsney-counts'
obama.saveAsTextFile(outputDir)
## @knitr map-reduce

### map-reduce step to sum hits across date-time-language triplets ###

def stratify(line):
    # create key-value pairs where:
    #   key = date-time-language
    #   value = number of website hits
    vals = line.split(' ')
    return(vals[0] + '-' + vals[1] + '-' + vals[2], int(vals[4]))

# sum number of hits for each date-time-language value
counts = obama.map(stratify).reduceByKey(add) # 5 minutes
# 128889 for full dataset

### map step to prepare output ###

def transform(vals):
    # split key info back into separate fields
    key = vals[0].split('-')
    return(",".join((key[0], key[1], key[2], str(vals[1]))))

### output to file ###

# have one partition because one file per partition is written out
mydir = '/global/scratch/esther730'
outputDir = mydir+'/' + 'Dinsney-counts'
counts.map(transform).repartition(1).saveAsTextFile(outputDir)

```

After above action, I got Disney data from the wikipages.

```

#copy the data from savio and rename it as disney_data
scp esther730@dtb.brc.berkeley.edu://global/scratch/esther730/Dinsney-counts2/part-00000 /mnt/c/Users

```

As showing in the plot, the highest hits appears three days before the Chirsmas.

```

# Note: refer from chapter 8 obama_plot.R
library(dplyr)
library(chron)

setwd("C:/Users/Esther/Desktop/stat243-fall-2017-master/section/s07")
dat <- read.table("disney_data", sep = ",")

```


```

names(dat) <- c("date", "time", "lang", "hits")
dat$date <- as.character(dat$date)
dat$time <- as.character(dat$time)
dat$time[dat$time %in% c("0", "1")] <- "000000"
wh <- which(nchar(dat$time) == 5)
dat$time[wh] <- paste0("0", dat$time[wh])
dat$chron <- chron(dat$date, dat$time, format = c(dates = "ymd", times = "hms"))
dat$chron <- dat$chron - 5/24 # GMT -> EST

dat <- dat %>% filter(dat$lang == "en")
# collect data between Chismas period in 2008. Bedtime Stories release in
# December 25, 2008.
sub <- dat %>% filter(dat$date < 20081231 & dat$date > 20081220)
plot(sub$chron, sub$hits, type = "l", xlab = "time", ylab = "hits", main = "Disney hits among Christmas
points(sub$chron, sub$hits)

```

Disney hits among Christmas period in 2008



```
# dev.off()
```

4 Q4

4.1 (a)

```
library(foreach)
require(parallel)
require(doParallel)
nCores <- as.integer(Sys.getenv("SLURM_CPUS_ON_NODE"))
registerDoParallel(nCores)
library(stringr)
library(readr)

#Note: The single function and usage of str_pad are refer from my classmate, Ming Qiu.
#write a function to detect the strings that contains "Barack_Obama"
single <- function(filename){
  data <- readLines(filename)
  library(stringr)
  out <- data[str_detect(data, 'Barack_Obama')]
  return(out) }

n <- 960
#nsub <- 10
result <- foreach(
  i = 1:n,
  .packages = c("stringr"),
  .combine = c,
  .verbose = TRUE) %dopar% {
  filename <- paste("/global/scratch/paciorek/wikistats_full/dated_for_R/part-", str_pad(i-1, width=5,
  #filename <- paste("C:/Users/Esther/Desktop/stat243-fall-2017-master/part-",str_pad(i-1, width=5,
  outputs <- single(filename)
  #outputs
}

#Transfer result into dataframe
results <- data.frame (matrix(sapply(list(result),function(x) unlist(strsplit(x,split=" ") ) ), ncol=6,

#results <- data.frame ( t(sapply(result,function(x) unlist(strsplit(x,split=" ") ) ) )
colnames(results) <- c("date", "time", "language", "webpage", "number of hits", "page size")
head(results)
class(results)
print(dim(results))
proc.time()
```

Part of output file is following:

```
dic <- "C:/Users/Esther/Desktop/stat243-fall-2017-master/section/s07/ps6q4no.out"
output <- readLines(dic)
tail(output, 25)

## [1] "      date   time  language                               webpage"
## [2] "1 20081231 100000 commons.m                               File:Barack_Obama_signature.svg"
```

```
## [3] "2 20081202 150000      en Image:Michelle,_Oprah_Winfrey_and_Barack_Obama.jpg"
## [4] "3 20081229 200001      en      Europe_and_Ireland_not_MYOB_or_Barack_Obama"
## [5] "4 20081123 100000      it      Discussion:Barack_Obama"
## [6] "5 20081101 010000      pt      Imagem:Barack_Obama.jpg"
## [7] "6 20081231 120000      en      Washington_Or_Barack_Obama"
## [8] "  number of hits page size"
## [9] "1          4      26796"
## [10] "2          2      19158"
## [11] "3          1       6047"
## [12] "4          1     40260"
## [13] "5          3     24840"
## [14] "6          1       6007"
## [15] "> class(results)"
## [16] "[1] \"data.frame\""
## [17] "> print(dim(results))"
## [18] "[1] 325263      6"
## [19] "> proc.time()"
## [20] "      user      system    elapsed "
## [21] "34664.559  2247.375  2715.079 "
## [22] "> "
## [23] "> proc.time()"
## [24] "      user      system    elapsed "
## [25] "34664.560  2247.375  2715.093 "
```

4.2 (b)

As the time is 2715 seconds(46 minutes) for one node, which vert close to 12 minutes if divide it my 4. Hence, using R for each has almost similar times or even faster.

4.3 (c)

```
#statistic allcate
library(foreach)
require(parallel)
require(doParallel)
nCores <- as.integer(Sys.getenv("SLURM_CPUS_ON_NODE"))
registerDoParallel(nCores)
library(stringr)

#Note: The single function and usage of str_pad are refer from my classmate, Ming Qiu.
single <- function(filename){
  data <- readLines(filename)
  library(stringr)
  out <- data[str_detect(data,'Barack_Obama')]
  return(out) }

n <- 960
#n <- 2
path <- "/global/scratch/paciorek/wikistats_full/dated_for_R/part-"
#path <- "C:/Users/Esther/Desktop/stat243-fall-2017-master/part-"
#obtains all the files' directories
i=1:n
filename <- mclapply(n, mc.preschedule=TRUE, function(x) paste(path,str_pad(i-1, width=5, side="left",
output <- mclapply(filename[[1]], single, mc.preschedule=TRUE)
```

```
#transfer the data type
results <- data.frame( matrix(sapply( unlist( output ), function(x) unlist(strsplit(x,split=" ") )), nc

colnames(results) <- c("date", "time", "language", "webpage", "number of hits", "page size")
head(results)
class(results)
print(dim(results))
proc.time()
write.table(results, file='/global/scratch/esther730/results_ob.csv',sep = ",",row.names = FALSE)
```

The following output reveal the processing time of static allocate, which is 17198 that nearly 289 minutes, when running at 4 cores, it will cost nearly 71 minutes.

```
dic <- "C:/Users/Esther/Desktop/stat243-fall-2017-master/section/s07/ps6q4st.out"
output <- readLines(dic)
tail(output, 25)

## [1] "1
## [2] "2 Special:AllPages/I_ran_Project_Vote_voter_registration_drive_in_Illinois,_ACORN_was_smack_da
## [3] "3
## [4] "4
## [5] "5
## [6] "6
## [7] " number of hits page size"
## [8] "1 86 2032215"
## [9] "2 2 25520"
## [10] "3 1 7825"
## [11] "4 16 760462"
## [12] "5 4 55875"
## [13] "6 1 20922"
## [14] "> class(results)"
## [15] "[1] \"data.frame\""
## [16] "> print(dim(results))"
## [17] "[1] 433895 6"
## [18] "> proc.time()"
## [19] " user system elapsed "
## [20] "16643.04 343.28 17197.31 "
## [21] "> write.table(results, file='/global/scratch/esther730/results_ob.csv',sep = \",\",row.names =
## [22] "> "
## [23] "> proc.time()"
## [24] " user system elapsed "
## [25] "16643.653 343.314 17197.955 "
```

```
#dynamic allocate
library(foreach)
require(parallel)
require(doParallel)
nCores <- as.integer(Sys.getenv("SLURM_CPUS_ON_NODE"))
registerDoParallel(nCores)
library(stringr)

single <- function(filename){
  data <- readLines(filename)
  library(stringr)
```



```

out <- data[str_detect(data,'Barack_Obama')]
return(out) }

n <- 960
#nsub <- 10
result <- foreach(
  i = 1:n,
  .packages = c("stringr"),
  .combine = c,
  .verbose = TRUE) %dopar% {
  filename <- paste("/global/scratch/paciorek/wikistats_full/dated_for_R/part-", str_pad(i-1, width=5,
  outputs <- mclapply(filename,single, mc.preschedule=FALSE) #false for dynamic
  #outputs
}
#transfer result into dataframe
results <- data.frame (matrix(sapply(list(result),function(x) unlist(strsplit(x,split=" ") ) ), ncol=6,
#results <- data.frame ( t(sapply(result,function(x) unlist(strsplit(x,split=" ") ) ) ) )
colnames(results) <- c("date", "time", "language", "webpage", "number of hits", "page size")
head(results)
class(results)
print(dim(results))
proc.time()

```

The dynamic allocate used 22585 seconds(377minutes) to grep the strings among 960 files. When running on four nodes, it will cost about 94 minutes to obtain the results.

```

dic <- "C:/Users/Esther/Desktop/stat243-fall-2017-master/section/s07/ps6q4dy.out"
output <- readLines(dic)
tail(output, 25)

## [1] "1
## [2] "2 Special:AllPages/I_ran_Project_Vote_voter_registration_drive_in_Illinois,_ACORN_was_smack_da
## [3] "3
## [4] "4
## [5] "5
## [6] "6
## [7] " number of hits page size"
## [8] "1      86    2032215"
## [9] "2       2     25520"
## [10] "3       1      7825"
## [11] "4      16    760462"
## [12] "5       4     55875"
## [13] "6       1     20922"
## [14] "> class(results)"
## [15] "[1] \"data.frame\""
## [16] "> print(dim(results))"
## [17] "[1] 433895      6"
## [18] "> proc.time()"
## [19] " user system elapsed "
## [20] "43138.56 1849.43 22584.29 "
## [21] "> #write.table(results, file='/global/scratch/esther730/results_ob.csv',sep = \",\\\",row.names
## [22] "> "
## [23] "> proc.time()"
## [24] " user system elapsed "
## [25] "43138.56 1849.43 22584.30 "

```

Dynamic allocation is better when iterations may take very different amounts of time. However, in this case, the data size for each file is quite similar, so it would be better to use static allocation which evenly divides loop iteration space into n chunks. Hence, generally, static allocation is a better method than dynamic way in this case.

5 Q5

5.1 (a)

(calculation attached as hand written version)

The cholesky will do $\frac{n^3}{6} + \frac{n^2}{2} - \frac{2n}{3}$ times operations.

5.2 (b)

Yes, we can overwrite the original matrix.

By the formula of cholesky decomposition, first, the a_{11} , which means the first element in the matrix, would only be used after the square root of it. Hence, it's fine to overwrite it.

Second, the first row will be calculated dividing each element to u_{11} , which is the first element of cholesky matrix we obtained by square root of a_{11} . Hence. According to the step 3 of the formula that begins to calculate the diagonal by row, and we found that only the elements after calculation will be used to obtain u_{ii} and u_{ij} (where i is smaller than j). Therefore, it's safe to overwrite the first row.

From the formula, we can find that when calculate u_{ii} and u_{ij} , what we need is the corresponding elements from original matrix and the transfer elements from the former rows. Hence, we won't use any corresponding elements of the former rows from the original matrix. Hence, we can overwrite the original matrix by row to each element.